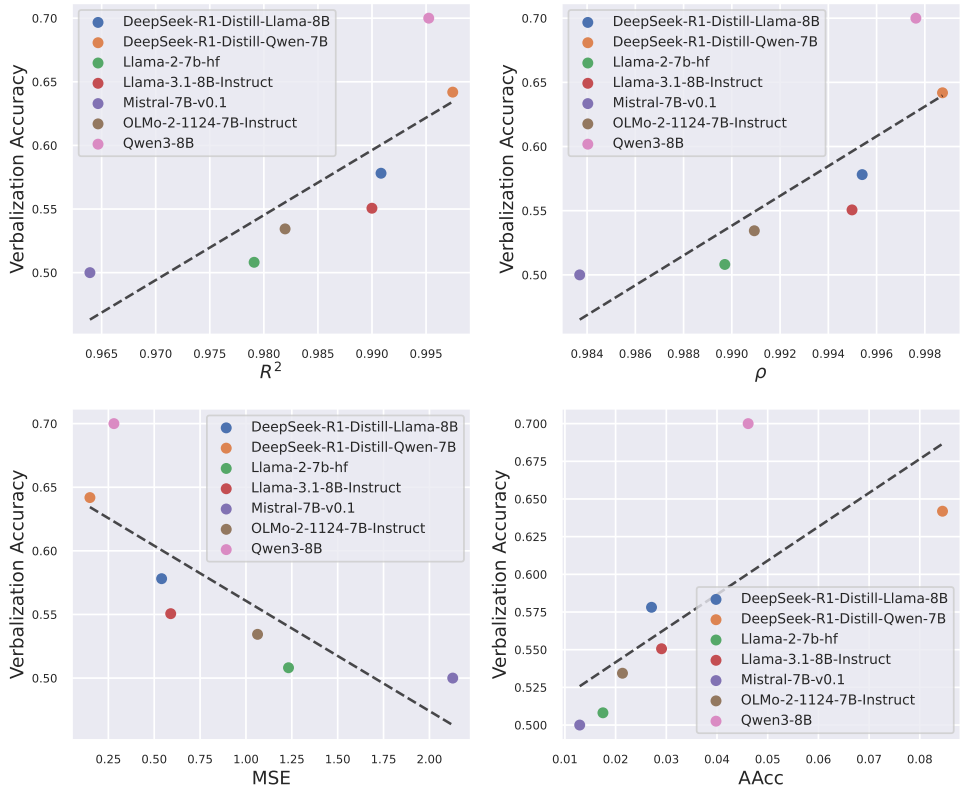

Figure 9: Average performance of linear regression probes at the first 3 layers correlates with model's verbalization accuracy on `cross-notation`. This is an expanded version of Fig. 5.

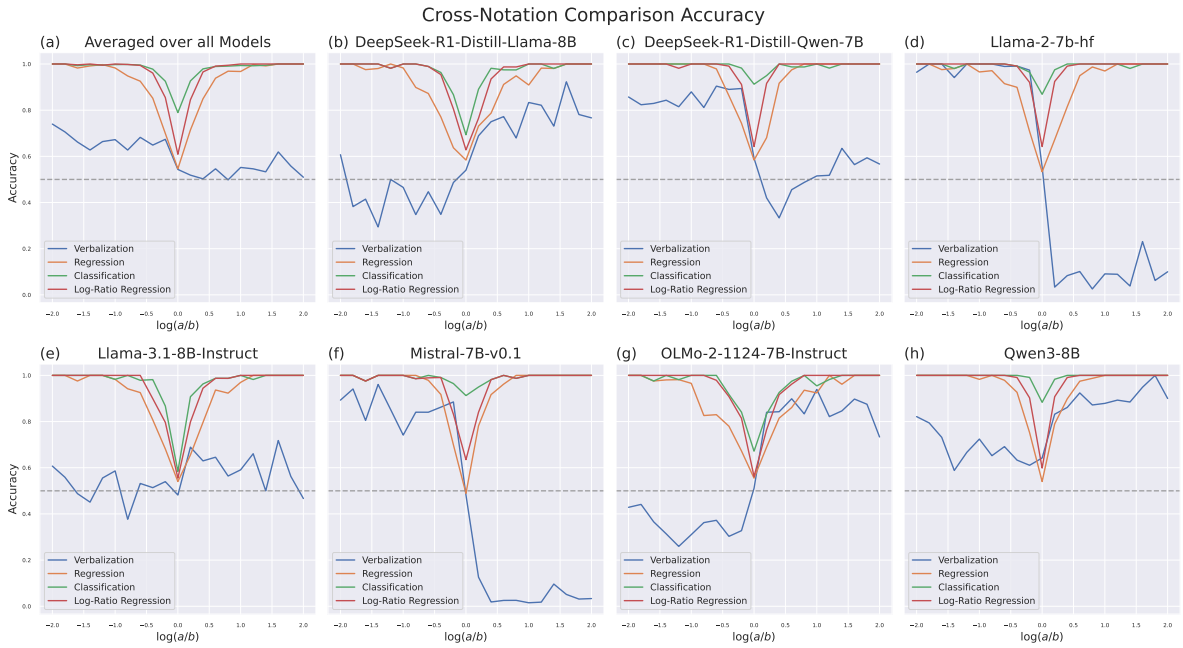

Figure 10: Cross-notation comparison accuracy with varying relative magnitude of two numbers (measured by $\log_2(a/b)$) using different methods. Same as Fig. 3, but with each models' results displayed separately.

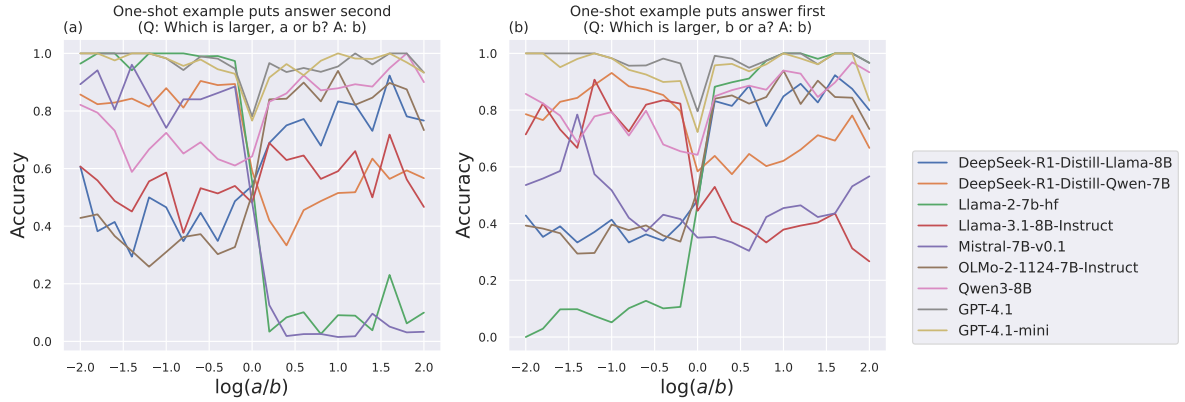

Figure 11: One-shot verbalized cross-notation comparison accuracy using different examples. (a) Same as Fig. 3(a); (b) Models' verbalized response when the numbers in the example are exchanged in positions.

Figure 12: Few-shot verbalized cross-notation comparison accuracy for different models. In each panel, color-matched dashed lines indicate the average accuracy of each shot count.

Figure 13: Cross-notation comparison accuracy vs. number of digits using different comparison methods. (a) Verbalized comparison using one-shot prompting; (b) Comparison via the predicted value of the regression probe trained on the hidden states; (c) Comparison using a logistic classification probe trained on the hidden states;

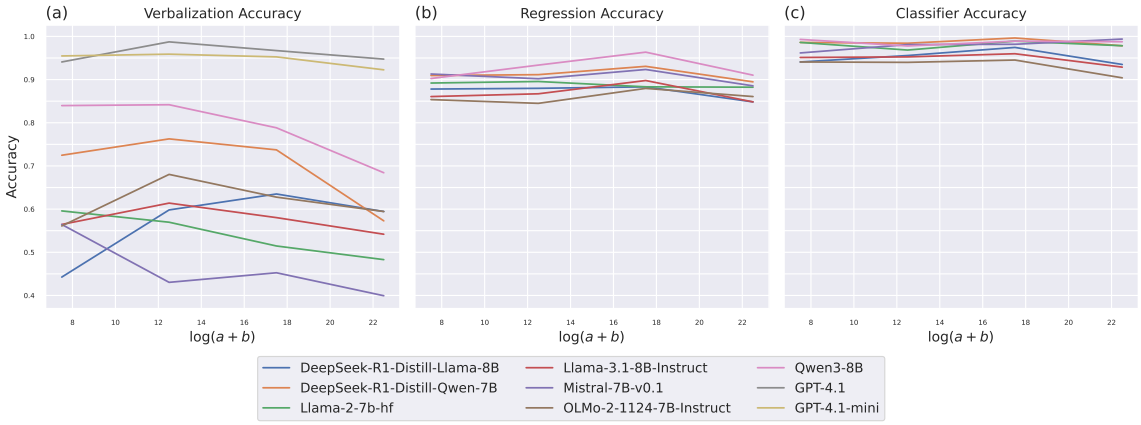

Figure 14: Cross-notation comparison accuracy vs. $\log_2(a+b)$ using different comparison methods. (a) Verbalized comparison using one-shot prompting; (b) Comparison via the predicted value of the regression probe trained on the hidden states; (c) Comparison using a logistic classification probe trained on the hidden states.